

ZFS is a rethinking of the storage stack, combining traditional file systems as well as volume managers into one cohesive tool. ZFS has some specific terminology that sets it apart from more traditional storage systems, however it has a great set of features with a focus on usability for systems administrators.

## ZFS Concepts

### Virtual Devices

A VDEV is similar to a raid device presented by a RAID card, there are several different types of VDEV's that offer various advantages, including redundancy and speed. In general VDEV's offer better reliability and safety than a RAID card. It is discouraged to use a RAID setup with ZFS, as ZFS expects to directly manage the underlying disks.

Types of VDEV's \* stripe (a single disk, no redundancy) \* mirror (n-way mirrors supported) \* raidz \* raidz1 (1-disk parity, similar to RAID 5) \* raidz2 (2-disk parity, similar to RAID 6) \* raidz3 (3-disk parity, no RAID analog) \* disk \* file (not recommended for production due to another filesystem adding unnecessary layering)

Your data is striped across all the VDEV's present in your Storage Pool, so more VDEV's will increase your IOPS.

### Storage Pools

ZFS uses Storage Pools as an abstraction over the lower level storage provider (VDEV), allow you to separate the user visible file system from the physical layout.

### ZFS Dataset

ZFS datasets are analogous to traditional filesystems but with many more features. They provide many of ZFS's advantages. Datasets support Copy on Write snapshots, quota's, compression and deduplication.

### Limits

One directory may contain up to  $2^{48}$  files, up to 16 exabytes each. A single storage pool can contain up to 256 zettabytes ( $2^{78}$ ) of space, and can be striped across  $2^{64}$  devices. A single host can have  $2^{64}$  storage pools. The limits are huge.

## Commands

### Storage Pools

Actions: \* List \* Status \* Destroy \* Get/Set properties

List zpools

```
# Create a raidz zpool
$ zpool create bucket raidz1 gpt/zfs0 gpt/zfs1 gpt/zfs2
```

```
# List ZPools
$ zpool list
NAME      SIZE  ALLOC   FREE  EXPANDSZ   FRAG    CAP  DEDUP  HEALTH  ALTROOT
```

```
zroot    141G    106G    35.2G          -    43%    75%    1.00x    ONLINE    -
```

*# List detailed information about a specific zpool*

```
$ zpool list -v zroot
```

NAME	SIZE	ALLOC	FREE	EXPANDSZ	FRAG	CAP	DEDUP	HEALTH	ALTROOT
zroot	141G	106G	35.2G	-	43%	75%	1.00x	ONLINE	-
gptid/c92a5ccf-a5bb-11e4-a77d-001b2172c655	141G	106G	35.2G		-	43%	75%		

Status of zpools

*# Get status information about zpools*

```
$ zpool status
```

```
pool: zroot
```

```
state: ONLINE
```

```
scan: scrub repaired 0 in 2h51m with 0 errors on Thu Oct 1 07:08:31 2015
```

```
config:
```

NAME	STATE	READ	WRITE	CKSUM
zroot	ONLINE	0	0	0
gptid/c92a5ccf-a5bb-11e4-a77d-001b2172c655	ONLINE	0	0	0

```
errors: No known data errors
```

*# Scrubbing a zpool to correct any errors*

```
$ zpool scrub zroot
```

```
$ zpool status -v zroot
```

```
pool: zroot
```

```
state: ONLINE
```

```
scan: scrub in progress since Thu Oct 15 16:59:14 2015
```

```
39.1M scanned out of 106G at 1.45M/s, 20h47m to go
```

```
0 repaired, 0.04% done
```

```
config:
```

NAME	STATE	READ	WRITE	CKSUM
zroot	ONLINE	0	0	0
gptid/c92a5ccf-a5bb-11e4-a77d-001b2172c655	ONLINE	0	0	0

```
errors: No known data errors
```

Properties of zpools

*# Getting properties from the pool properties can be user set or system provided.*

```
$ zpool get all zroot
```

NAME	PROPERTY	VALUE	SOURCE
zroot	size	141G	-
zroot	capacity	75%	-
zroot	altroot	-	default
zroot	health	ONLINE	-
...			

*# Setting a zpool property*

```
$ zpool set comment="Storage of mah stuff" zroot
```

```
$ zpool get comment
NAME    PROPERTY  VALUE                SOURCE
tank    comment   -                    default
zroot    comment   Storage of mah stuff local
```

Remove zpool

```
$ zpool destroy test
```

## Datasets

Actions: \* Create \* List \* Rename \* Delete \* Get/Set properties

Create datasets

```
# Create dataset
$ zfs create tank/root/data
$ mount | grep data
tank/root/data on /data (zfs, local, nfsv4acls)

# Create child dataset
$ zfs create tank/root/data/stuff
$ mount | grep data
tank/root/data on /data (zfs, local, nfsv4acls)
tank/root/data/stuff on /data/stuff (zfs, local, nfsv4acls)
```

```
# Create Volume
$ zfs create -V zroot/win_vm
$ zfs list zroot/win_vm
NAME                USED  AVAIL  REFER  MOUNTPOINT
tank/win_vm          4.13G  17.9G   64K    -
```

List datasets

```
# List all datasets
$ zfs list
NAME                USED  AVAIL  REFER  MOUNTPOINT
zroot               106G  30.8G   144K   none
zroot/ROOT          18.5G  30.8G   144K   none
zroot/ROOT/10.1       8K    30.8G   9.63G   /
zroot/ROOT/default   18.5G  30.8G   11.2G   /
zroot/backup         5.23G  30.8G   144K   none
zroot/home           288K   30.8G   144K   none
...
```

```
# List a specific dataset
$ zfs list zroot/home
NAME                USED  AVAIL  REFER  MOUNTPOINT
zroot/home          288K  30.8G   144K   none
```

```
# List snapshots
```

```
$ zfs list -t snapshot
zroot@daily-2015-10-15                                0      -    144K
zroot/ROOT@daily-2015-10-15                           0      -    144K
zroot/ROOT/default@daily-2015-10-15                   0      -   24.2G
zroot/tmp@daily-2015-10-15                             124K   -    708M
zroot/usr@daily-2015-10-15                             0      -    144K
zroot/home@daily-2015-10-15                           0      -   11.9G
zroot/var@daily-2015-10-15                           704K   -   1.42G
zroot/var/log@daily-2015-10-15                       192K   -    828K
zroot/var/tmp@daily-2015-10-15                        0      -    152K
```

Rename datasets

```
$ zfs rename tank/root/home tank/root/old_home
$ zfs rename tank/root/new_home tank/root/home
```

Delete dataset

```
# Datasets cannot be deleted if they have any snapshots
zfs destroy tank/root/home
```

Get / set properties of a dataset

```
# Get all properties
```

```
$ zfs get all zroot/usr/home
NAME                PROPERTY            VALUE                SOURCE
zroot/home          type                filesystem            -
zroot/home          creation            Mon Oct 20 14:44 2014 -
zroot/home          used                11.9G                -
zroot/home          available            94.1G                -
zroot/home          referenced            11.9G                -
zroot/home          mounted            yes                  -
...
```

```
# Get property from dataset
```

```
$ zfs get compression zroot/usr/home
NAME                PROPERTY            VALUE                SOURCE
zroot/home          compression          off                  default
```

```
# Set property on dataset
```

```
$ zfs set compression=zip-9 mypool/lamb
```

```
# Get a set of properties from all datasets
```

```
$ zfs list -o name,quota,reservation
NAME                QUOTA    RESERV
zroot               none     none
zroot/ROOT          none     none
zroot/ROOT/default  none     none
zroot/tmp           none     none
zroot/usr           none     none
zroot/home          none     none
zroot/var           none     none
...
```

## Snapshots

ZFS snapshots are one of the things about zfs that are a really big deal

- The space they take up is equal to the difference in data between the filesystem and its snapshot
- Creation time is only seconds
- Recovery is as fast as you can write data.
- They are easy to automate.

Actions: \* Create \* Delete \* Rename \* Access snapshots \* Send / Receive \* Clone

Create snapshots

```
“bash # Create a snapshot of a single dataset zfs snapshot tank/home/sarlalian@now
```

## Create a snapshot of a dataset and its children

```
$ zfs snapshot -r tank/home@now $ zfs list -t snapshot NAME USED AVAIL REFER MOUNTPOINT
tank/home@now 0 - 26K - tank/home/sarlalian@now 0 - 259M - tank/home/alice@now 0 - 156M -
tank/home/bob@now 0 - 156M - ...
```

Destroy snapshots

```
# How to destroy a snapshot
```

```
$ zfs destroy tank/home/sarlalian@now
```

```
# Delete a snapshot on a parent dataset and its children
```

```
$ zfs destroy -r tank/home/sarlalian@now
```

Renaming Snapshots

```
# Rename a snapshot
```

```
$ zfs rename tank/home/sarlalian@now tank/home/sarlalian@today
```

```
$ zfs rename tank/home/sarlalian@now today
```

```
# zfs rename -r tank/home@now @yesterday
```

Accessing snapshots

```
# CD Into a snapshot directory
```

```
$ cd /home/.zfs/snapshot/
```

Sending and Receiving

```
# Backup a snapshot to a file
```

```
$ zfs send tank/home/sarlalian@now | gzip > backup_file.gz
```

```
# Send a snapshot to another dataset
```

```
$ zfs send tank/home/sarlalian@now | zfs recv backups/home/sarlalian
```

```
# Send a snapshot to a remote host
```

```
$ zfs send tank/home/sarlalian@now | ssh root@backup_server 'zfs recv tank/home/sarlalian'
```

```
# Send full dataset with snapshots to new host
```

```
$ zfs send -v -R tank/home@now | ssh root@backup_server 'zfs recv tank/home'
```

## Cloning Snapshots

```
# Clone a snapshot
$ zfs clone tank/home/sarlalian@now tank/home/sarlalian_new

# Promoting the clone so it is no longer dependent on the snapshot
$ zfs promote tank/home/sarlalian_new
```

## Putting it all together

This following a script utilizing FreeBSD, jails and ZFS to automate provisioning a clean copy of a mysql staging database from a live replication slave.

```
#!/bin/sh

echo "==== Stopping the staging database server ====="
jail -r staging

echo "==== Cleaning up existing staging server and snapshot ====="
zfs destroy -r zroot/jails/staging
zfs destroy zroot/jails/slave@staging

echo "==== Quiescing the slave database ====="
echo "FLUSH TABLES WITH READ LOCK;" | /usr/local/bin/mysql -u root -pmyrootpassword -h slave

echo "==== Snapshotting the slave db filesystem as zroot/jails/slave@staging ====="
zfs snapshot zroot/jails/slave@staging

echo "==== Starting the slave database server ====="
jail -c slave

echo "==== Cloning the slave snapshot to the staging server ====="
zfs clone zroot/jails/slave@staging zroot/jails/staging

echo "==== Installing the staging mysql config ====="
mv /jails/staging/usr/local/etc/my.cnf /jails/staging/usr/local/etc/my.cnf.slave
cp /jails/staging/usr/local/etc/my.cnf.staging /jails/staging/usr/local/etc/my.cnf

echo "==== Setting up the staging rc.conf file ====="
mv /jails/staging/etc/rc.conf.local /jails/staging/etc/rc.conf.slave
mv /jails/staging/etc/rc.conf.staging /jails/staging/etc/rc.conf.local

echo "==== Starting the staging db server ====="
jail -c staging

echo "==== Make sthe staging database not pull from the master ====="
echo "STOP SLAVE;" | /usr/local/bin/mysql -u root -pmyrootpassword -h staging
echo "RESET SLAVE;" | /usr/local/bin/mysql -u root -pmyrootpassword -h staging
```

## Additional Reading

- BSDNow's Crash Course on ZFS

- [FreeBSD Handbook on ZFS](#)
- [BSDNow's Crash Course on ZFS](#)
- [Oracle's Tuning Guide](#)
- [OpenZFS Tuning Guide](#)
- [FreeBSD ZFS Tuning Guide](#)